Executive Summary

Chewy's CPFR program is transforming how supply-chain data is shared and understood. Today, weekly emailed reports—each built from multiple queries—and parallel VDS processes, fragment visibility, duplicate effort, and limit scalability.

This initiative unites two complementary components: a centralized data platform in Snowflake that delivers governed, daily CPFR data at scale, and a CPFR Vendor Portal within Chewy Partner Hub (CPH), that offers vendors secure, always-on access through a modern interface. Together, these capabilities create a shared, consistent foundation that strengthens vendor collaboration and focus, simplifies access, and positions the organization for sustained collaborative growth. The initiative delivers $1.4M in recurring annual CPFR-specific operational value (ISM capacity expansion, data governance efficiency, cost avoidance enablement), with additional unmeasured vendor-side efficiency gains strengthening collaborative relationships. Combined with VC and Chargeback integration, this unified platform investment supports Chewy's digital modernization strategy while establishing infrastructure for future advanced analytics capabilities.

Context and Current State

Across Chewy's supply chain, CPFR activities span multiple processes and workstreams that depend on parallel data pulls, manual coordination, and frequent reconciliation. These distributed practices create friction that limits scale and consistency.

Today, CPFR distributes weekly inventory and forecast reports to more than 2,300 vendors. While some steps are automated, each cycle still depends on semi-manual initiation, recurring adjustments to contact and report configurations, and frequent re-sends. These repetitive tasks—along with the corrective work that follows minor data discrepancies—inflate otherwise avoidable effort and consume time that could be redirected toward proactive vendor engagement. This cycle also introduces the risk of version drift, when teams or vendors reference slightly different report iterations and data.

VDS currently provides on-demand dashboards for a smaller premium vendor group. Its session-based model is effective for deep exploration and security isolation, yet scales poorly to thousands of concurrent vendors and complicates historical comparisons. In parallel, the Vendor Compliance (VC) team and Replen Tech are developing a portal on CPH. This presents a timely opportunity to align CPFR modernization with an existing, funded initiative. Through co-development, CPFR contributes governance and data design, while leveraging Replen Tech's product management, engineering, and EDS security resources.

Entitlement Overview — Classes of Benefit

1. Operational Cohesion and Efficiency

A unified platform eliminates redundant action and manual reconciliations. Internal teams and vendors will draw from the same governed data sources, ensuring consistent interpretation of fill rate, PDP, forecast, and other key metrics. Automated refresh and role-based access that ensures each vendor sees only their own data replace brittle report assembly for both outbound reports and internal analyses. Survey evidence documents 1-5 hours per ISM/week spent on manual data prep and error correction—capacity that will be redirected to planning, root-cause analysis, and strategic vendor engagement.

2. Vendor Enablement and Partnership Expansion

Within CPH, vendors gain continuous, self-service access to accurate daily data filtered to their specific SKUs. They will also be able to update contact details and communication preferences directly—streamlining simple administrative tasks without changing oversight or control by ISMs or leadership. This model supports thousands of concurrent vendor tenants and provides a scalable path toward future API connectivity. The result is stronger collaboration, faster issue resolution, and smoother onboarding of new vendor partners—reducing friction so ISMs can focus on substantive planning conversations rather than routine data delivery.

3. Enterprise Data Alignment and Future Readiness

The platform establishes a single governed source of CPFR data shared across CPFR, VC, Merch, and related teams. Standardized definitions and lineage reduce planning risk and eliminate conflicting extracts. Daily immutable snapshots ensure consistency, auditability, and historical traceability. To support additions or updates to collaborative data, changes can be made at the single source, flowing immediately and uniformly to all users. The structured design also lays the groundwork

for advanced analytics, machine learning, and AI-driven automation, minimizing the preparation traditionally required. What began as operational reporting becomes a durable enterprise asset supporting Chewy's digital evolution.

4. Strategic Growth Enablement

By combining broad vendor access with governed data consistency, Chewy extends data-driven collaboration across its entire vendor network while maintaining VDS's role for advanced analytics. The same architecture can readily expand to include Vendor Compliance (VC) and related datasets, providing a unified environment for vendor-specific information. This shared platform reduces duplication, simplifies vendor engagement, and establishes a foundation for future cross-program dashboards and enterprise-wide visibility.

Quantitative Snapshot

Benefit Category	Estimation Basis	Appx. 1st-Yr Value	Key Assumption(s)
Operational Cohesion & Efficiency	36 team members × 2 avg hours/week reallocated to higher-value work	**$ 608 k**	Capacity redirected from manual prep to analysis
Vendor Enablement & Partnership Expansion	Faster, more frequent and more actionable vendor contribution	**$ 458 k**	Faster vendor decisions, reduced friction
Enterprise Data Alignment & Future Readiness	Fewer data inconsistencies, rework, or avoidable costs	**$ 339 k**	Improved accuracy across teams
Strategic Growth Enablement	*Synergy benefits and shared infrastructure reuse*	***$ 625 k****	*Multi-program applicability*
	Total Conservative Estimate	***≈ $ 1.4 M per year***	*Recurring post-stabilization*

** Not included in current entitlement claim. This represents value of road-mapped future phases of expansion*

Dependencies, Risks & Mitigations

Risk / Dependency	Mitigation
Data consistency across systems	Centralized governance and shared change-management process
Vendor adoption of portal	Co-branded rollout with VC and Replen Tech; vendor education and support
Inter-team alignment	Defined data charters and joint implementation roadmap
Platform sustainment	Ownership by Supply-Chain B.I. with ongoing support from Replen Tech and EDS

Entitlement Benefit

The CPFR Platform and Portal initiative advances Chewy's goals of operational excellence and partner enablement by transforming manual reporting into a governed, always-available data capability. In partnership with VC and Chargeback, it enables consistent, data-driven collaboration across connected supply-chain functions. The initiative yields an estimated $1.4 million in recurring annual benefit specific to CPFR—building on and amplifying the value generated by VC and Chargeback integration. This combined momentum directly supports Chewy's digital modernization strategy and represents a pragmatic, low-risk investment in sustained growth and data confidence.

Appendix 1: Valuation Basis

The valuation basis reflects conservative estimates of CPFR's indirect contribution to operational efficiency and cost avoidance. As a collaborative planning enabler, CPFR's value accrues through improved data visibility, reduced coordination friction, and enhanced analytical capacity—factors that support broader organizational effectiveness but do not generate direct revenue.

[1] Operational Efficiency Calculations

Factor	Value	Notes
ISM cohort	36	Includes Consumables, Health/Science, EFE, Hard Goods,
Hours per person/week	2	Median surveyed response value
Rate	64.90	C5 average (likely higher IRL)
Weeks	52	Burden rate does not stop during PTO/Holidays
FBLR	2.5	Inflation factor for fringe, overhead, G&A--based on GAO guidelines and ACEC/AEC project estimators for white-collar roles
Annual Op Ex	607500	

[2] Vendor Enablement & Partnership Expansion

Factor	Value	Notes
Vendor friction reduction	0.5	Avg hours saved per week via self-service
Vendor opportunity cost rate	50	Conservative estimate of fungible vendor contribution recovery
Chewy capture rate	.15	Attribution factor for Chewy enablement
Annual Op Ex	*458250*	

[3] Enterprise Data Alignment

Factor	Value	Notes
Annual Major COE Events	12	Inventory efficiency gains, COE, reduced spoilage, NOP improvement
Cost Avoidance per event	415000	Average of high-impact events
Annual Large NPI Events	10	NPI/Catalog additions and product line unlocks.
Incremental SKU revenue	45000	Average across In-Stock
Impact rate of opportunity cost + COE	0.0625	Avg impact of CPFR activities – derived from weekly task proportion
Annual Op Ex	339375	

[4] Strategic Growth Enablement

Factor	Value	Notes
Analytic Opportunity	150000	Pro rata value of program contribution: future analytics opportunities
AI/ML Opportunity	375000	Pro rata value of program contribution: future AI/ML opportunities
Lean / 6σ	100000	Pro rata value of program contribution : future Lean / 6σ opportunities
Annual Cap Ex Conversion	625000	

Appendix 2: Case Studies

Survey Methodology & Scope

Subject matter expert interviews were conducted with senior In-Stock Managers managing relationships with major vendors and direct distributors. The ISMs surveyed collectively represent substantial CPFR program scope spanning a substantial portion of COGS and revenue. Interviews focused on quantifying operational burden, identifying data consistency challenges, and documenting vendor feedback regarding access preferences and collaborative planning friction.

1. Manual Coordination Overhead & Operational Burden

Time Investment in Data Preparation and Validation

ISMs report consistent weekly time burdens associated with CPFR data preparation, validation, and correction of errors (COE):

- **1-2 hours per week**: Typical for relationships with stable data flows and mature current-state CPFR processes
- **2-3 hours per week**: Common where improved consistency and reduced COE would yield measurable impact
- **5+ hours per week**: Reported in cases where email-based data exchange creates unnecessary churn and where CPFR consistency, availability, or COE issues compound

These time investments represent capacity diverted from strategic vendor collaboration, root-cause analysis, and proactive planning work.

End-of-Period Escalations

Transportation and MOQ mismatch issues create predictable escalation patterns. ISMs report spending up to **2 days out of certain weeks** rectifying MOQ and transportation constraints that surface at period boundaries (month-end, quarter-end). While improving visibility around these issues extends beyond the initial CPFR platform scope, the frequency and impact of these escalations underscore the value of incremental feature that can be added once foundational portal infrastructure is in place.

Discontinued SKU Management

Substantial ISM time is consumed rectifying discontinued SKU status and cleanup. This represents a recurring risk of pricing mismatches and operational friction that could be streamlined through improved data visibility and self-service capabilities.

2. Data Visibility & Consistency Challenges

Non-Optimal Placement (NOP) Visibility

NOP data requires enhanced visibility with daily (or more frequent) cadence to effectively track when corrections have been absorbed into planning cycles. Current weekly cadence creates delays in understanding whether corrective actions have taken effect, extending the feedback loop between issue identification validation, and further actions.

Future Availability Date (FAD) Lists

FAD lists signal to ordering automation when to pause orders until a specified future date. These require broader visibility—preferably real-time—to ensure updates are reflected immediately. Current fragmented visibility creates coordination gaps between vendor manufacturing schedules and Chewy ordering systems.

New Item Forecasting

New product introductions often require off-cadence forecast updates and more frequent refreshes. Typical NPI forecasting takes multiple cycles to stabilize, creating early-stage planning friction. While CPFR data access improvements may not directly accelerate NPI forecast generation, platform capabilities can support faster distribution once forecasts are available.

Forecast Mismatch Thresholds & Promotional Tensions

Forecast accuracy thresholds create significant financial exposure when misaligned with promotional sales goals. Threshold misses directly impact vendor payback guarantees, creating incentives that can drive inefficient inventory decisions.

Past instances of threshold-driven forecast mismatch have resulted in approximately **$5 million in unwarranted buy-up** due to unclear visibility into forecast accuracy relative to promotional expectations. Based on analysis of recurring threshold miss patterns, this issue class represents an estimated **$15 million annually** in avoidable cost if not corrected through improved data visibility and governance.

The structural tension between promotional incentives (which increase buy-up pressure) and in-stock efficiency goals (which prioritize inventory optimization) creates recurring exposure. Major vendor relationships experience this dynamic at least annually, making improved forecast visibility and consistent data sharing critical to reducing financial waste.

3. Vendor Self-Service Preferences & API Demand

Direct Vendor Feedback: API Access Preference

In conference calls, **Blue Buffalo** expressed distinct preference for API access to CPFR data, explicitly weighing this capability against the cost of VDS tenancy (3% of COGS up to $400K plus $250K). Blue Buffalo characterized current email-based CPFR data transmission as a "struggle" and indicated that API availability would be a significant decision factor in their vendor data services investment evaluation.

This preference echoes previous statements from **Mars** and **Nestle-Purina**—the two largest vendors in the CPFR program. The consistency of this feedback across major vendor relationships indicates that API access is not an isolated request but a systematic preference among sophisticated vendor partners, and much simpler to deploy via CPH portal.

Portal Access Prioritization

Vendor feedback emphasizes portal access over advanced thresholding or signaling features for initial rollout. Self-service data access is viewed as delivering greater immediate value in both vendor-Chewy alignment and collaborative follow-up than incremental analytical sophistication (though this was a close second). This suggests a clear implementation priority: establish reliable self-service access first, then expand analytical capabilities based on adoption patterns.

Daily Data Cadence Requirements

Vendors have explicitly requested **daily-available data** via portal, particularly for:

- PDP out-of-stock (OOS) metrics
- Fill rate performance
- Forecast updates

Recovery cycles for PDP and fill rate issues often occur over days rather than weeks, making weekly data cadence insufficient for responsive corrective action. Daily availability is estimated to save approximately **1 additional hour per week** in coordination overhead per affected vendor, while enabling up to a **50% improvement in OOS recovery** for affected items.

While affected items represent about 0.1% of company-wide OOS, affected vendors—by nature—trend higher than average, suggesting the impacted inventory pool may approach 1% for the vendors most likely to use daily data for rapid response.

4. SKU and Vendor Number Assignment Cleanup

Visibility into unused-but-not-deactivated vendor numbers and agreement lines would facilitate cleanup of a particularly insidious data quality issue: price discrepancies between active and obsolete vendor numbers for the same SKU.

Example scenario: A vendor's primary vendor number receives a price increase (reflected in the active agreement). However, an obsolete vendor number tied to the same SKU—with a lower historical price—remains active but does not receive price updates. Pricing algorithms select the lower (invalid) price, creating potential discrepancies.

This issue is especially acute because:

- Vendors all too often do not know their own vendor numbers
- Vendor numbers are not **easily** searchable
- Portal ticket submissions lack adequate fields for vendor number entry
- Long-term vendors often accumulate multiple active agreements and vendor numbers, plus expired/inactive entries

Improved visibility would help prevent items from being added to incorrect vendor numbers and enable systematic cleanup of legacy data structures.

Summary of Quantified Impacts

The survey evidence documents measurable operational burden and financial exposure:

Time burden reduction opportunities**:**

- 1-5 hours per ISM per week in CPFR data preparation and COE
- ~2 days per week addressing MOQ/transportation escalations (opportunity for incremental feature expansion)
- ~1 hour per week per vendor relationship through daily data availability

Financial impact from improved visibility:

- $5M documented cost avoidance case (single threshold mismatch event)
- ~$15M estimated annual exposure from recurring forecast threshold issues

Vendor relationship evidence:

- Consistent API access preference across major vendors (Blue Buffalo, Mars, Nestle-Purina)
- Explicit daily data cadence requirements for operational responsiveness
- Clear prioritization of portal self-service over advanced analytical features

This evidence base supports the business case for CPFR platform modernization by documenting specific, quantifiable pain points experienced by senior ISMs managing substantial vendor relationship portfolios. The convergence of operational burden, financial exposure, and vendor demand creates a compelling case for investment in governed data infrastructure and self-service portal capabilities.

Appendix 3: Architectural Summary (Platform)

- Illustration is representational-only.
- Segregating Portal access and Chewy TM access allows clear separation of access scopes while maintaining synchronization to single-source data.

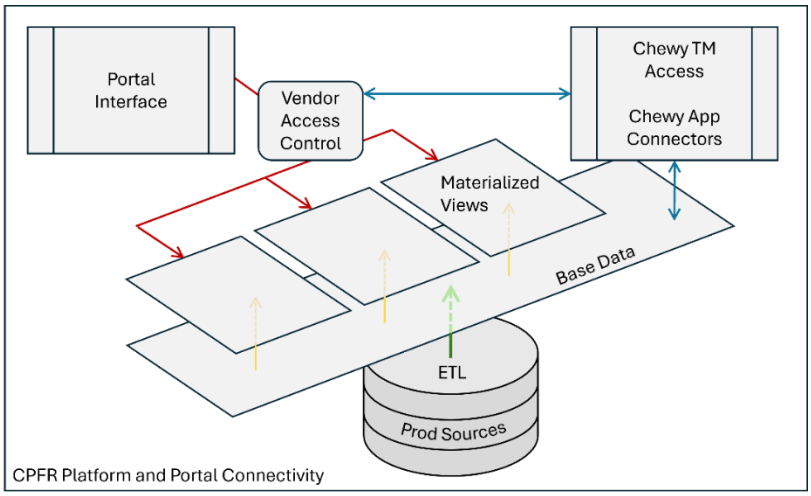

CPFR Platform and Portal Connectivity

Appendix 4: Implementation Plan

Build & Rollout Strategy:

*The rollout will follow a **phased approach** that balances speed with control.*

- **Phase 1** establishes the **data platform foundation** in Snowflake and deliver internal tools for In-Stock and B.I. teams.
- **Phase 2** will extend capabilities to the **CPFR Vendor Portal** within Chewy Partner Hub (CPH), enabling secure vendor access and self-service.

While these phases are sequential, a degree of overlap between late-stage Phase 1 and early-stage Phase 2 is expected and even desirable. This approach minimizes disruption and allows technical, training, and communication activities to proceed in parallel as readiness milestones are met.

Additional measures to ensure a smooth transition:

- In-Stock staff will receive pre-release training and reference documentation prior to portal enablement.
- Early rollout will focus on a smaller initial vendor group (e.g. Tier 1 vendors followed by Tier 2, etc.), allowing Chewy to validate workflows and resolve any early issues before broad release.
- Lessons learned from this initial onboarding will inform subsequent vendor activation waves.

Communication Strategy:

Throughout build and rollout, progress and reference materials will be tracked in a shared collaboration space (e.g., Confluence). Communication to internal and external users will be staged to align with the two primary phases:

- Pre-launch notifications: sent 30 days prior to each rollout milestone, specifying go-live timing and hypercare period.
- Orientation sessions: offered for internal users and early vendor adopters during the hypercare window (≈ 30 days post-launch). Additional sessions may be added based on adoption needs.
- Awareness updates: periodic “coming soon” highlights during development to prepare teams for upcoming functionality.

Hypercare:

- A defined period of elevated service attention will follow each rollout milestone to ensure stable adoption:
- Externally, hypercare maintains Chewy’s commitment to partnership by providing responsive support and encouraging early adoption success.
- Internally, it delivers fast feedback loops for support and engineering teams to implement quick fixes and optimizations.

Email Backup:

An email feed capability will be available for vendors unable or with inopportune access to the portal. This will be activated at portal rollout to send the same data-view shown to the vendor in the portal based on the metric configuration chosen by the vendor’s ISM and selections the vendor makes in the portal for their recipients. Legacy email automation will be available as a back-up facility for CPFR reporting during rollout, but will be gradually deactivated in favor of portal-based systems.