CPFR Data-Sharing: Current State & Candidate Alternatives

Chewy's vendor-facing data-shares currently operate through two primary channels: CPFR and VDS. CPFR serves as the foundational data-sharing platform, reaching approximately 70% of Chewy's vendor population through weekly emailed reports containing tactical inventory and forecast data. In contrast, VDS operates as a premium, revenue-generating service offering strategic insights to a select group of vendors through dashboard and API access. As a new candidate system for data shares, CPH can host vendor access if data pipelines or platforms (collectively, "services") have already been established.

	CPFR	VDS Premium	VDS Direct	CPH
Purpose	Provides fundamental data used by vendors and ISMs for Inventory and order planning	Provide a broader set of data to facilitate insights and deep analysis of supply, operations, and performance	Provides highly granular and customizable data feeds that directly integrate with vendor systems	Aggregates Services from across Chewy, providing a single point of access for vendors
Share Method	Emailed	Dashboard (Tableau Server)	API access	Custom portal
Level of Granularity	SKU/weekly	SKU/daily; SKU/weekly	SKU/daily; SKU/weekly (depending on subscription)	TBD
Principle Characteristics	• Weekly 'snapshot' data • "Tactical" data-scope • Principal performance metrics • Tiered subsets of data • Fixed aggregation levels	• On-demand access • "Strategic" data-scope • Weekly or daily refreshed, depending on type and source • Limited ability to 'drill-down' into data	• On-demand access • Weekly or daily refreshes • "Exploratory" data-scope • Permits complex queries • Direct vendor system integration	• On-demand access • Customizable views and data content
Data-Shared	• Forecast • Autoship • Inventory • OOS • Fill Rate • DOS • NOP • Catalog Details ***CPFR Knowledge Domain*** ***VDS Knowledge Domain***	• Autoship Analysis • Inventory & Forecast • Location Out of Stock • PDP Out of Stock • Procurement Scorecard • Customer Experience Insights • Attachment Detail • Basket Overview • Competitive Comparison • Returns • Sales & Geographical Metrics • Trial & Repeat	[Daily & Weekly]: • Sales • Autoship • Receipt / Returns • Trial Rate / Review Avg. • ASN & Compliance • Inventory direct feed • Forecast direct feed	Same as CPFR *CPH Provides a portal—data pipelines or platform must be provided by Business Intelligence or Software Engineers (or both)*
Revenue Model	None (Free)	3% of COGS, $400K max	+ $100K (weekly); +$250K (daily)	TBD (Presumably None/Free)
Users Served	• ≈1,650 FC-inventory vendors • ≈625 Drop-ship vendors	• 65 unique vendors • Includes Top 20	• 5 unique vendors • All within Top 10	• ≈1,650 FC-inventory vendors • ≈625 Drop-ship vendors
Use-Cases	• "Plan-of-record" forecasting • Joint business planning • Vendor scorecard reference • Tracking inventory health • over/under performance alerts	• Root-cause analysis (vendor) • Sales auditing • EDI issue confirmation • Consumer purchase behavior • Competitor analysis	• Complex/granular analysis • POS • Direct database integration	• Same as CPFR

CPFR Data-Sharing: Current State & Candidate Alternatives

APPENDIX 1: CPH Implementation

To consider CPH as a potential sharing solution for CPFR data, we must first recognize that CPH functions as a conduit and an aggregator for nearly any kind of service that Chewy wishes to attach to it, however it doesn't automatically supply the underlying data services. In CPFR's case, this means that the service connecting source data to the hub would need to be provided and configured as an additional workstream. This creates two areas of development that must be coordinated. The first being the CPH hub presence itself, and the second is the data service supplying the hub.

	CPH Hub + Raw CPFR Queries	CPH Hub + data-service	CPH Hub + BI-built data-service	CPH Hub + SWE/BI data-service
Jobs to be Done	CPH builds the hub portal, and uses CPFR-supplied queries to pipe as-is from Snowflake	CPH builds portal, and also lead data-service development using Software Engineers (SWEs) working in consultation with CPFR team	CPH builds portal, Supply Chain Business Intelligence Engineering teams (SC-BIEs) build and deliver production-read data-service in consultation with CPFR team	CPH builds portal, SC-BIE + CPFR prototype data-service, and SWEs work in consultation with the above teams to integrate for production
CPFR Team LOE *(Level of effort)*	Low at build High to sustain	Mid at build Mid to sustain	Mid to build (work is front-loaded) Low to sustain	Mid to build (work is front-loaded) Low to sustain
CPH Team LOE	Mid at build Low to sustain	High at build Low to sustain	Mid at build Low to sustain	Mid at build Low to sustain
Software Engineering Team(s) LOE	Low at build n/a—sustainment	High at build Mid to sustain	Low at build n/a—sustainment	Mid at build Low to sustain
Business Intelligence Team(s) LOE	Low at build Mid to sustain	n/a—build Low to sustain	High at build Low to sustain	High at build Low to sustain
Data Governance Capability	Low / None	Good—possibly slow to update	High	High
Clearly Separates CPFR / VDS Knowledge Domains	No	Good—may *drift* if updates are sporadic	High	High
Enhanced Data Availability	No, and tiers are static	Yes, tiers may be slow to update, limited schema changes	Yes, full schema control	Yes, full schema control
Can Release MVP Platform Before Hub is Complete	No	No	Yes	Maybe
Advantages	Fastest CPH roll-out	Single threaded developers	Single-threaded sustainment; "in-house" enhancement capabilities	"In-house" enhancement capabilities
Complications / Risks	Inherits any current suboptimalities or errors in queries; limited extensibility	Multi-threaded sustainment; externalizes design, build, and future enhancements	Must coordinate two build teams in separate domains	Must coordinate three build teams in overlapping domains

CPFR Data-Sharing: Current State & Candidate Alternatives

APPENDIX 2: CPFR Tiering Structure

CPFR's tiered data access structure demonstrates how Chewy balances comprehensive vendor coverage with differentiated data needs. The four-tier system serves 2,274 total vendors, with Tier 1 (10 vendors) receiving the most comprehensive dataset including extended forecasting horizons, detailed inventory metrics, and complete catalog information. As tiers progress from 1 to 4, data granularity decreases while vendor count increases significantly—Tier 4 alone serves 1,255 vendors with basic operational metrics. This structure reflects Chewy's approach to scaling data-sharing across a diverse vendor base, with ISMs tailoring data access to align with each vendor's operational capabilities and business needs. This illustrates the depth and breadth of CPFR data that services must support.

CPFR	Tier 4	Tier 3	Tier 2	Tier 1
Forecast	Current Month Forecast Next 3-Months Forecast	Current Month Forecast Next 6-Months Forecast	Current Month Forecast Next 6-Months Forecast	Current Month Forecast Next 6-Months Forecast
Autoship			AS Demand (30 days)	AS Demand (30 days) AS-Backorders
Inventory			OH Units OO Units	OH Units OO Units
OOS			PDP % (trailing 30-day avg)	PDP % (trailing 30-day avg)
Fill Rate			Fill-Rate % (trailing 30-day avg)	Fill-Rate % (trailing 30-day avg)
DOS			DOS	DOS
NOP				NOP by FC NOP by Region NOP / OP Demand Total Demand
Catalog Details				Published Y/N Discontinued Y/N MOQ Base UOM, Purchase UOM Eaches per Case/Layer/Pallet Order divisibility by Pallet/Layer Temp Disable

Vendors Per Tier *	1255	955	54	10

** Includes FC + Drop-ship vendors; Tiers 2 - 4 are sent by unique vendor number… Tier 1 considers only parent company*

*** CPFR enrolls approximately 70% of overall vendor population*

CPFR Data-Sharing: Current State & Candidate Alternatives

APPENDIX 3: Notes and FAQ

Additional Considerations:

1. Buy vs. Build: At this time, Chewy's CPFR needs are best served by improving consistency and accessibility of data. Chewy's tech-stack and in-house expertise is already competent for these purposes. 'Off the shelf' tools or platforms are not expected to add significant incremental value and potentially increase risk/cost.
2. VDS Lite: While VDS Lite should be considered a relevant candidate among these solutions, VDS's present status in KTLO, as well as the undefined nature of the 'Lite' offering's capacities and capabilities make it difficult to accurately evaluate in comparison to the other options listed here. If VDS Lite requires further build-out, development might be protracted. However, if more substantial appraisals of VDS Lite's capabilities, capacity and readiness become available, it's suitability should be reconsidered. Until that time, CPFR solutions should not be predicated on VDS Lite.

Frequently Asked Questions:

- ***Can CPFR be migrated to CPH now—without waiting on data-service build-out***
  Yes, but this doesn't solve data quality issues, and would actually make it more difficult to isolate and fix errors or misalignment in sources.
- ***Why aren't we giving all data across VDS, CPFR (including its different tiers), and VC to vendors?***
  The reasons for segregating data-shares are due to observed differences in the analytical capabilities of different vendors (think: Nestle vs Mom & Pop), the way that "law of small numbers" can skew perceptions of trend for small vendors, and the responsibility for managing data-driven adverse-actions among a large number of vendors using a relatively smaller In-stock staff. However, reduction of data-sharing segments might further streamline operations or open-up opportunities for both vendors and Chewy, so this topic merits periodic reevaluation.
- ***Why is an external portal better than the emails we currently use for CPFR?***
  In addition to immediate reductions in CPFR administrative overhead, a single point-of-access vastly improves the consistency of experience for vendors, and simplifies future efforts to improve accuracy, consistency, and cadence of data delivery.
- ***Why won't a portal or external presence solve the internal needs of Chewy In-stock as well, when asking about the same data?***
  While it wouldn't be a problem for Chewy Team Members to use the a portal for "at a glance" data viewing, Chewy TM's frequently need to combine data for deep-dives, and may have good reasons to substitute alternate data definitions at times (think: different forecast models, fill-rate calculations, etc.).
- ***What is meant by data- "pipeline", "platform", or collectively "service"? ...does this require new business software?***
  These terms all refer to data management tools or assets like data tables and queries that "sit" between primary database sources and the final viewer tools (Tableau, et al) or portals. "Pipelines" and "platforms" differ in their approach, and "data-service" is a catch-all term for both—but in all cases, these assets seek to pre-aggregate data using clear definitions that avoid inconsistent or errant representations of business data. A well-governed data-service also vastly improves the speed of troubleshooting, root-cause analysis, or making tactical changes to CPFR data (think: changing a KPI calculation). Chewy already employs all the business / I.T. tools needed to support CPFR's data-service needs.
- ***What's the user going to see in a portal?***
  The portal can be configured in different ways, but for CPFR purposes the most important thing to showcase is the "last mile" view of the data, after data-services have transformed and compiled it. This will resemble the tables sent in emails, but with more advanced ability to filter, and summarize.
- ***What is the Chewy Team Member's experience accessing data in these different cases?***
  Chewy Team Members should retain their ability to access data independently from the portal. However, the use-cases for CPFR present an unusually acute need for well-defined, and consistent data sources. With this in mind, Chewy TMs would benefit from data-service assets as well, albeit with broader access.