
Executive Summary

Chewy's CPFR program is modernizing how supply-chain data is shared and interpreted. Today, CPFR information travels through weekly emailed reports assembled from multiple queries or via parallel processes in VDS—a process that fragments visibility, duplicates effort, and limits scalability.

The proposed initiative introduces two complementary components. A **centralized data platform** in Snowflake will provide governed, high-scale access to consistent, daily CPFR data. A **CPFR Vendor Portal** within Chewy Partner Hub (CPH), built in partnership with the Replen Tech team, will deliver secure, 24/7 vendor access through a modern interface. Together these elements create a shared foundation that unifies Chewy and vendor perspectives, simplifies access, and positions the organization for future analytical growth.

The combined benefit potential is estimated at **≈ \$1.5 – 1.7 million per year** in recurring value through workflow efficiency, reduced reconciliation effort, faster vendor alignment, and expansion of vendor participation. Beyond direct savings, the initiative establishes CPFR's data platform as a core enabler for cross-team collaboration and future AI-driven insight.

Context and Current State

CPFR currently distributes weekly inventory and forecast reports to more than 2,300 vendors. Though components of the process are automated, each cycle requires manual coordination of data extraction, validation, and delivery. This approach constrains scale, introduces risk of version drift, and consumes Team Member time better used for proactive vendor support.

VDS also serves CPFR data to a smaller premium vendor group through on-demand dashboards. Its session-based design isolates each query run—an effective model for deep exploration (not to mention, infosec)—but one that scales poorly to thousands of concurrent vendors and complicates time-series comparisons.

Meanwhile, the Vendor Compliance (VC) team and Replen Tech are building a portal on CPH, presenting a timely opportunity to align CPFR's modernization with an existing, funded framework. Co-development allows CPFR to contribute governance and data structure while leveraging Replen Tech's product-management, data-engineering, and EDS security resources.

Entitlement Overview — Classes of Benefit

1. Operational Cohesion & Efficiency

A unified data platform eliminates redundant queries and reconciliations. Internal users and vendors will rely on the same governed tables, ensuring consistent interpretation of fill-rate, PDP, forecast accuracy, and other key metrics. Automated refresh and isolation controls replace manual report assembly, freeing capacity for ISMs and analysts to focus on higher-value planning, root-cause exploration, and vendor collaboration. The effect is equivalent to adding new analytical planning staff without actually adding new headcount, redirecting expertise from routine reporting to insight generation.

2. Vendor Enablement & Partnership Expansion

Within CPH, vendors will gain continuous, on-demand access to consistent daily data, filtered to their specific SKUs. They will manage contacts and notifications directly, reducing administrative cycles for both the CPFR and In-Stock teams. This model supports thousands of concurrent vendor tenants, and provides a natural path for future API connectivity. The result is stronger collaboration, quicker issue resolution, and a scalable mechanism to engage new vendors.

3. Enterprise Data Alignment & Future Readiness

The platform component creates a single governed source of CPFR data accessible to CPFR, VC, Merch, and other teams. Standardized definitions and lineage eliminate the planning risk that arises from conflicting extracts. Daily immutable snapshots maintain consistency and auditability while enabling historical comparison. The structured design also prepares CPFR data for advanced analytics, machine learning, and other automation tools—transforming what was once operational reporting into a reliable enterprise asset.



42 **4. Strategic Growth Enablement**

43 By combining broad vendor access with governed data consistency, Chewy extends data-driven collaboration across its full
44 vendor base while preserving VDS’s role for premium analytics. The same architecture can be augmented to store and manage
45 Vendor Compliance (VC) data, offering a unified environment for vendor-specific information. This shared platform reduces
46 duplication, simplifies vendor interactions, and lays the groundwork for future shared dashboards and cross-program visibility.

47 **[WIP] Quantitative Snapshot**

Benefit Category	Estimation Basis	Appx. 1st-Yr Value	Key Assumption(s)
Operational Cohesion & Efficiency	20 team members × 4 hours/week reallocated to higher-value work	\$ 608 k	Capacity redirected from manual prep to analysis
Vendor Enablement & Partnership Expansion	2,350 vendors × avg \$ 300 value / vendor	\$ 712 k	Faster vendor decisions, reduced friction
Enterprise Data Alignment & Future Readiness	Fewer data inconsistencies & rework	\$ 375 k	Improved accuracy across teams
Strategic Growth Enablement	Synergy benefits and shared infrastructure reuse	\$ 625 k*	Multi-program applicability
Total Conservative Estimate		≈ \$ 1.7 M per year	Recurring post-stabilization

48 * Not included in current entitlement claim. This represents value of road-mapped future phases of expansion

49 **Dependencies, Risks & Mitigations**

Risk / Dependency	Mitigation
Data consistency across systems	Centralized governance and shared change-management process
Vendor adoption of portal	Co-branded rollout with VC and Replen Tech; vendor education and support
Inter-team alignment	Defined data charters and joint implementation roadmap
Platform sustainment	Ownership by Supply-Chain B.I. with ongoing support from Replen Tech and EDS

50 **Entitlement Benefit**

51 The CPFR Platform and Portal initiative advances Chewy’s goals of operational excellence and partner enablement by
52 converting manual, report-based workflows into a governed, always-available data capability. CPFR’s platform and portal
53 initiative warrants OP1/OP2 support for joint development with VC and Replen Tech, with an expected recurring benefit of
54 approximately \$1.7 million per year and clear alignment to Chewy’s digital modernization strategy, this initiative represents a
55 pragmatic, low-risk investment in sustained growth and data confidence.

58 **[WIP] Appendix:**

- 59
- **Valuation Basis:** Summary of entitlement basis and projected Chewy benefit
 - **Case Studies:** Examples of quantitative or qualitative opportunities, with fiscal and/or operational basis and detail.
 - **Architecture Summary:** Snowflake data platform → governed views → CPH portal.
 - **Implementation Plan:** Phase 1 – Platform build & internal tools; Phase 2 – Portal integration & vendor onboarding.
 - **Next Steps:** Coordinate with B.I. and Replen Tech to confirm ownership and implementation roles. Confirm approval.
- 64

Appendix 1: Valuation Basis

[1] Operational Efficiency Calculations

Factor	Value	Notes
ISM Cohort	36	Includes Consumables, Health/Science, EFE, Hard Goods,
Hours per person/week	2	Survey sample of cohort and adjust
Rate	64.90	C5 average (likely higher IRL)
Weeks	52	Burden rate does not stop during PTO/Holidays
FBLR	2.5	Inflation factor for fringe, overhead, G&A--based on GAO guidelines and ACEC/AEC project estimators for white-collar roles
Annual Op Ex	607500	

[2] Vendor Enablement & Partnership Expansion

Factor	Value	Notes
Vendor Cohort	2350	Enrolled CPFR tenants only
Avg Annual Sales	10000	Compute from cohort
COE incremental lift	0.01	Survey sample of cohort and adjust
Percent encumbered	0.03	Survey ISMs and adjust
Annual Op Ex	712050	

[3] Enterprise Data Alignment & Future Readiness

Factor	Value	Notes
Annual Deep-dive Events (Lift)	30	Annual incremental revenue from throughput improvements
Lift per event	45000	Average across In-Stock
Annual Deep-dive Events (COE)	60	Operational efficiency gains, COE, Spoilage avoidance, NOP improvement
Cost Avoidance per event	45000	Average across In-Stock
Annual Deep-dive Events (NPI)	10	NPI/Catalog additions and product line unlocks.
Incremental SKU revenue	45000	Average across In-Stock
Impact rate of data inconsistency	0.1	Percent of time that more data consistency or visibility was impactful
Annual Op Ex	450000	

[4] Strategic Growth Enablement

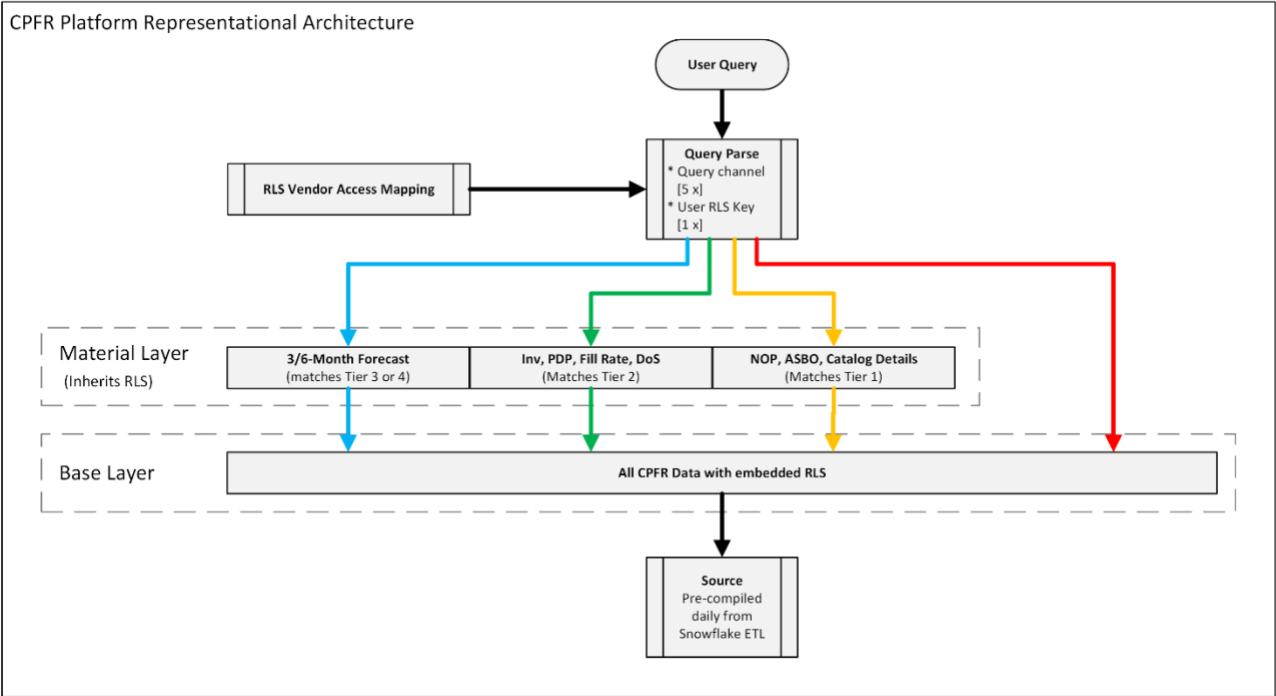
Factor	Value	Notes
Analytic Opportunity	150000	Pro rata value of program contribution: future analytics opportunities
AI/ML Opportunity	375000	Pro rata value of program contribution: future AI/ML opportunities
Lean / 6σ	100000	Pro rata value of program contribution : future Lean / 6σ opportunities
Annual Cap Ex Conversion	625000	

68
69
70
71


Appendix 2: Case Studies

[WIP—placeholder]

DRAFT



- The illustration above is representational only. Colored paths above loosely represent different request types.
- Platform resolves (refreshes) all data at the same time, **once-daily**
- The red line above illustrates direct access to base layer data, which may be relevant for:
 - Specific, business-critical augments (e.g. Vendor 1 depends heavily on an additional metric which is not useful to other vendors). This ensures that load and frequency for critical 1-off requests does not impact other vendors. This also ensures that any unique fields are still scoped and governed in the overall data dictionary and policy
 - Permits direct tool access for Chewy internal users to make cross-vendor, or widely-scoped queries without direct impact on external vendor-facing service levels



Portal and Platform Connectivity

- Illustration above is representational.
- Segregating Portal access and Chewy TM access allows clear separation of access scopes while maintaining synchronization to single-source data.

Appendix 4: Implementation Plan

Build & Rollout Strategy:

The rollout will follow a **phased approach** that balances speed with control.

- **Phase 1** establishes the **data platform foundation** in Snowflake and deliver internal tools for In-Stock and B.I. teams.
- **Phase 2** will extend capabilities to the **CPFR Vendor Portal** within Chewy Partner Hub (CPH), enabling secure vendor access and self-service.

While these phases are sequential, a degree of overlap between late-stage Phase 1 and early-stage Phase 2 is expected and even desirable. This approach minimizes disruption and allows technical, training, and communication activities to proceed in parallel as readiness milestones are met.

Additional measures to ensure a smooth transition:

- In-Stock staff will receive pre-release training and reference documentation prior to portal enablement.
- Early rollout will focus on a smaller initial vendor group (e.g. Tier 1 vendors followed by Tier 2, etc.), allowing Chewy to validate workflows and resolve any early issues before broad release.
- Lessons learned from this initial onboarding will inform subsequent vendor activation waves.

Communication Strategy:

Throughout build and rollout, progress and reference materials will be tracked in a shared collaboration space (e.g., Confluence). Communication to internal and external users will be staged to align with the two primary phases:

- **Pre-launch notifications:** sent 30 days prior to each rollout milestone, specifying go-live timing and the duration of the hypercare period.
- **Orientation sessions:** offered for internal users and early vendor adopters during the hypercare window ($\approx$ 30 days post-launch). Additional sessions may be added based on adoption needs.
- **Awareness updates:** periodic “coming soon” highlights during development to prepare vendors and internal teams for upcoming functionality.

Hypercare:

A defined period of elevated service attention will follow each rollout milestone to ensure stable adoption.

- Externally, hypercare maintains Chewy’s commitment to partnership by providing responsive support and encouraging early adoption success.
- Internally, it delivers fast feedback loops that allow support and engineering teams to implement quick fixes and optimizations.

Email Backup:

An email feed capability will be available for vendors unable or with inopportune access to the portal. This will be activated at portal rollout to send the same data-view shown to the vendor in the portal based on the metric configuration chosen by the vendor’s ISM and selections the vendor makes in the portal for their recipients. Legacy email automation will be available as a back-up facility for CPFR reporting during rollout, but will be gradually deactivated in favor of portal-based systems.